

## ✓ CNNEx1\_CatsDogs

Source: <https://www.kaggle.com/c/dogs-vs-cats/data>

```
!wget --no-check-certificate \
  https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip \
  -O /tmp/cats_and_dogs_filtered.zip

--2023-12-18 05:46:32-- https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.212.207, 142.250.128.207, 172.253.114.207, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.212.207|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 68606236 (65M) [application/zip]
Saving to: '/tmp/cats_and_dogs_filtered.zip'

/tmp/cats_and_dogs_100%[=====>] 65.43M  187MB/s   in 0.3s

2023-12-18 05:46:33 (187 MB/s) - '/tmp/cats_and_dogs_filtered.zip' saved [68606236/68606236]
```

```
#extract files from zip to /tmp
import os
import zipfile

local_zip = '/tmp/cats_and_dogs_filtered.zip'

zip_ref = zipfile.ZipFile(local_zip, 'r')

zip_ref.extractall('/tmp')
zip_ref.close()
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
#create necessary directories to split data
base_dir = '/tmp/cats_and_dogs_filtered'

train_dir = os.path.join(base_dir, 'train')
test_dir = os.path.join(base_dir, 'validation')

# train directories
train_cats_dir = os.path.join(train_dir, 'cats')
train_dogs_dir = os.path.join(train_dir, 'dogs')

# test directories
test_cats_dir = os.path.join(test_dir, 'cats')
test_dogs_dir = os.path.join(test_dir, 'dogs')
```

```
#Check data shape
print('Training (Cat) :', len(os.listdir(train_cats_dir)))
print('Training (Dog) :', len(os.listdir(train_dogs_dir)))


print('Testing (Cat) :', len(os.listdir(test_cats_dir)))
print('Testing (Dog) :', len(os.listdir(test_dogs_dir)))
#Expected output
#1000
#1000
#500
#500

Training (Cat) : 1000
Training (Dog) : 1000
Testing (Cat) : 500
Testing (Dog) : 500
```

```
#See the cats and dogs
%matplotlib inline

import matplotlib.image as mpimg
import matplotlib.pyplot as plt
#img_path=os.path.join(train_cats_dir, 'cat.40.jpg')
img_path=os.path.join(train_dogs_dir, 'dog.90.jpg')
img = mpimg.imread(img_path)
plt.imshow(img)
```

<matplotlib.image.AxesImage at 0x7d35a4161a20>



```
#Generate the data on the fly
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# reshape the image
train_datagen = ImageDataGenerator( rescale = 1.0/255. )
test_datagen = ImageDataGenerator( rescale = 1.0/255. )

#batch_size = 20 and target_size = 150x150
train_generator = train_datagen.flow_from_directory(train_dir,
                                                    batch_size=20,
                                                    class_mode='binary',
                                                    target_size=(150, 150))

test_generator = test_datagen.flow_from_directory(test_dir,
                                                  batch_size=20,
                                                  class_mode = 'binary',
                                                  target_size = (150, 150))
```

Found 2000 images belonging to 2 classes.  
Found 1000 images belonging to 2 classes.

```
#define the model
import tensorflow as tf
model = tf.keras.models.Sequential([
    #Conv2D adds a convolution layer, 16: number of filters(https://lodev.org/cgtutor/filtering.html), (3,3) convolution
    tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(150, 150, 3)), #images are 150x150 RGB(3)
    tf.keras.layers.MaxPooling2D(2,2), #MaxPooling2D will reduce the dimension
    tf.keras.layers.Conv2D(32, (3,3), activation='relu'), #another conv layer with 32 filters
    tf.keras.layers.MaxPooling2D(2,2), #another pooling layer
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'), #another conv layer with 64 filters
    tf.keras.layers.MaxPooling2D(2,2), #another pooling layers
    # Flatten the results to feed into a DNN
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'), #512 neurons
    tf.keras.layers.Dense(1, activation='sigmoid') #0 for cats and 1 for dogs (Binary classifier)
])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d (MaxPooling2D)	(None, 74, 74, 16)	0

conv2d_1 (Conv2D)	(None, 72, 72, 32)	4640
max_pooling2d_1 (MaxPoolin g2D)	(None, 36, 36, 32)	0
conv2d_2 (Conv2D)	(None, 34, 34, 64)	18496
max_pooling2d_2 (MaxPoolin g2D)	(None, 17, 17, 64)	0
flatten (Flatten)	(None, 18496)	0
dense (Dense)	(None, 512)	9470464
dense_1 (Dense)	(None, 1)	513

=====  
Total params: 9494561 (36.22 MB)  
Trainable params: 9494561 (36.22 MB)  
Non-trainable params: 0 (0.00 Byte)

```
#compile with RMSprop: RMS = Root Mean Squared  
from tensorflow.keras.optimizers import RMSprop
```

```
model.compile(optimizer=RMSprop(learning_rate=0.001),  
              loss='binary_crossentropy',  
              metrics = ['accuracy'])
```

```
#Fit the model with 15 epochs
```

```
history = model.fit(train_generator,  
                    validation_data=test_generator,  
                    steps_per_epoch=100,  
                    epochs=15,  
                    validation_steps=50,  
                    verbose=2)
```

```
Epoch 1/15  
100/100 - 5s - loss: 9.7066e-07 - accuracy: 1.0000 - val_loss: 2.7263 - val_accuracy: 0.7410 - 5s/epoch - 45ms/step  
Epoch 2/15  
100/100 - 6s - loss: 8.6776e-07 - accuracy: 1.0000 - val_loss: 2.7479 - val_accuracy: 0.7410 - 6s/epoch - 61ms/step  
Epoch 3/15  
100/100 - 4s - loss: 7.8723e-07 - accuracy: 1.0000 - val_loss: 2.7655 - val_accuracy: 0.7430 - 4s/epoch - 45ms/step  
Epoch 4/15  
100/100 - 4s - loss: 7.2243e-07 - accuracy: 1.0000 - val_loss: 2.7792 - val_accuracy: 0.7430 - 4s/epoch - 45ms/step  
Epoch 5/15  
100/100 - 6s - loss: 6.6388e-07 - accuracy: 1.0000 - val_loss: 2.7947 - val_accuracy: 0.7430 - 6s/epoch - 61ms/step  
Epoch 6/15  
100/100 - 4s - loss: 6.1564e-07 - accuracy: 1.0000 - val_loss: 2.8087 - val_accuracy: 0.7430 - 4s/epoch - 45ms/step  
Epoch 7/15  
100/100 - 9s - loss: 5.8168e-07 - accuracy: 1.0000 - val_loss: 2.8193 - val_accuracy: 0.7430 - 9s/epoch - 92ms/step  
Epoch 8/15  
100/100 - 5s - loss: 5.4558e-07 - accuracy: 1.0000 - val_loss: 2.8311 - val_accuracy: 0.7400 - 5s/epoch - 45ms/step  
Epoch 9/15  
100/100 - 6s - loss: 5.1336e-07 - accuracy: 1.0000 - val_loss: 2.8411 - val_accuracy: 0.7390 - 6s/epoch - 60ms/step  
Epoch 10/15  
100/100 - 4s - loss: 4.7973e-07 - accuracy: 1.0000 - val_loss: 2.8551 - val_accuracy: 0.7430 - 4s/epoch - 45ms/step  
Epoch 11/15  
100/100 - 6s - loss: 4.6241e-07 - accuracy: 1.0000 - val_loss: 2.8623 - val_accuracy: 0.7390 - 6s/epoch - 61ms/step  
Epoch 12/15  
100/100 - 4s - loss: 4.4097e-07 - accuracy: 1.0000 - val_loss: 2.8706 - val_accuracy: 0.7390 - 4s/epoch - 44ms/step  
Epoch 13/15  
100/100 - 4s - loss: 4.2050e-07 - accuracy: 1.0000 - val_loss: 2.8817 - val_accuracy: 0.7390 - 4s/epoch - 45ms/step  
Epoch 14/15  
100/100 - 6s - loss: 4.0507e-07 - accuracy: 1.0000 - val_loss: 2.8901 - val_accuracy: 0.7390 - 6s/epoch - 61ms/step  
Epoch 15/15  
100/100 - 5s - loss: 3.8955e-07 - accuracy: 1.0000 - val_loss: 2.8989 - val_accuracy: 0.7390 - 5s/epoch - 45ms/step
```

```

#The internal process
import numpy as np
import random
from tensorflow.keras.preprocessing.image import img_to_array, load_img

successive_outputs = [layer.output for layer in model.layers[1:]]

visualization_model = tf.keras.models.Model(inputs = model.input, outputs = successive_outputs)

cat_img_files = [os.path.join(train_cats_dir, f) for f in ['cat.0.jpg', 'cat.10.jpg', 'cat.25.jpg', 'cat.40.jpg',
                                                            'cat.100.jpg']]
dog_img_files = [os.path.join(train_dogs_dir, f) for f in ['dog.0.jpg', 'dog.10.jpg', 'dog.25.jpg', 'dog.40.jpg',
                                                            'dog.100.jpg']]

img_path = random.choice(cat_img_files + dog_img_files)
img = load_img(img_path, target_size=(150, 150))

x = img_to_array(img)
x = x.reshape((1,) + x.shape)

x /= 255.0

successive_feature_maps = visualization_model.predict(x)

layer_names = [layer.name for layer in model.layers]

for layer_name, feature_map in zip(layer_names, successive_feature_maps):

    if len(feature_map.shape) == 4:

        n_features = feature_map.shape[-1]
        size       = feature_map.shape[ 1]

        display_grid = np.zeros((size, size * n_features))

        for i in range(n_features):
            x = feature_map[0, :, :, i]
            x -= x.mean()
            x /= x.std ()
            x *= 64
            x += 128
            x = np.clip(x, 0, 255).astype('uint8')
            display_grid[:, i * size : (i + 1) * size] = x # Tile each filter into a horizontal grid

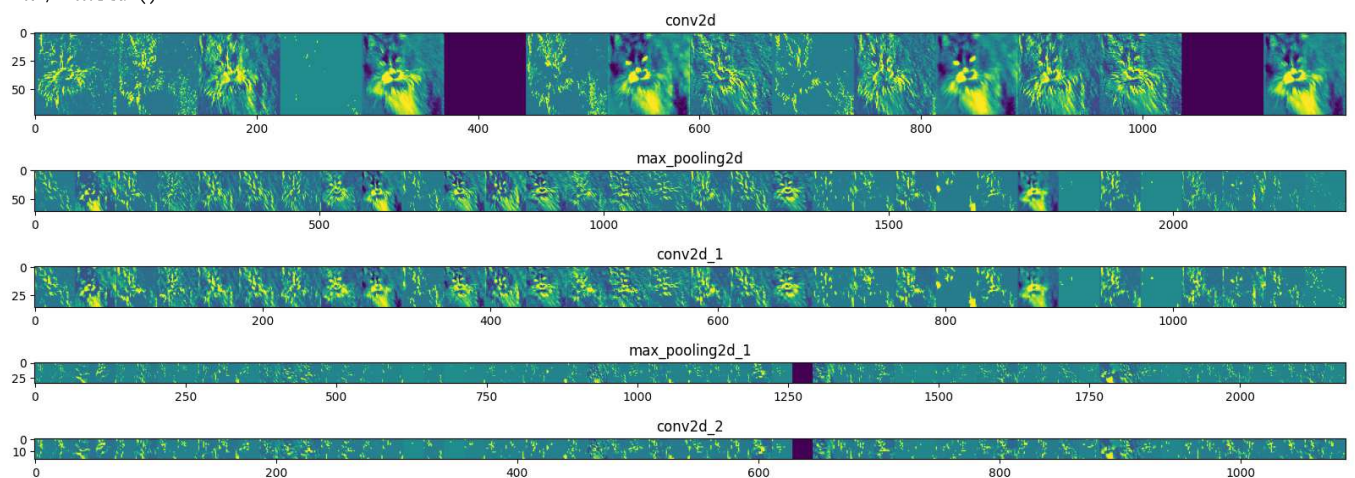
        scale = 20. / n_features
        plt.figure( figsize=(scale * n_features, scale) )
        plt.title ( layer_name )
        plt.grid ( False )
        plt.imshow( display_grid, aspect='auto', cmap='viridis' )

```

```

1/1 [=====] - 0s 307ms/step
<ipython-input-11-5fdf8221a19d>:41: RuntimeWarning: invalid value encountered in divide
x /= x.std ()

```



```
#Let's test with random images
import numpy as np

from google.colab import files
from keras.preprocessing import image

uploaded = files.upload()

for fn in uploaded.keys():
    path='/content/' + fn      #Save the image to content folder
    img=image.load_img(path, target_size=(150, 150))    #load the image

    x=image.img_to_array(img)
    x=np.expand_dims(x, axis=0)
    images = np.vstack([x])

    classes = model.predict(images, batch_size=10)    #predict the label for the image

    print(classes[0])    #Print the label, remember it will be either one or zero

    if classes[0]>0:
        print(fn + " is a dog")    #print human readable label

    else:
        print(fn + " is a cat")    #print human readable label
```

No file chosen  
enable.

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to